Towards Evaluation of Multi-party Dialogue Systems

Authors:

Khyati Mahajan ✉ kmahaja2@uncc.edu
Sashank Santhanam
Samira Shaikh

Main Motivation

- Prolific research in NLG evaluation
  - Multiple taxonomies presented[1, 2, 3, 4]
  - Studies towards importance of automatic and human metrics[5, 6, 7, 8]
  - + Confusion surrounding inconsistent evaluation methods used[9]
- However, not much work towards evaluation specifically for Multi-party Conversation (MPC) evaluation
  - = Need for discussing MPC specific challenges and needs

MPC Challenges

- The presence of multiple participants introduces new and interesting challenges from a dialogue modeling perspective
  - Participant roles - need to maintain speaker-specific and addressee-specific information jointly with dialogue modeling
  - Conversation structure - more graph-like than sequential
  - Threads within conversation - multiple topic threads could co-exist within sub-groups

Contributions

- Propose an expanded taxonomy focusing on the specific challenges introduced by multi-party dialogue, or group conversations
  - Such as the need to maintain speaker-specific context and recognize the proper addressees
- Synthesize evaluation measures utilized in existing MPC research, and relate them to the expanded taxonomy introduced
  - Report important inconsistencies in current research

Expanded Taxonomy

	Violation of Form	Violation of Content
Utterance	(I1) Uninterpretable (I2) Grammatical error	(I3) Semantic error (I4) Wrong information
Response	(I5) Ignore question (I6) Ignore request (I7) Ignore proposal (I8) Ignore greeting	(I9) Ignore expectation *(I18) Forgot speaker* *(I19) Forgot addressee(s)*
Context	(I10) Unclear intention (I11) Topic transition error (I12) Lack of information	(I13) Self-contradiction (I14) Contradiction (I15) Repetition
Society	(I16) Lack of sociality	(I17) Lack of common sense
Participant	*(I20) Wrong speaker* *(I21) Wrong addressee(s)*	*(I22) Wrong thread response* *(I23) Inappropriately timed initiative*

Table 1: Integrated taxonomy for errors in chat-oriented dialogue systems by Higashinaka et al. (2021). We extend the taxonomy to include errors specific to MPD - extensions are italicized and highlighted in grey. The numbering is assigned serially and used in text to refer to discussions surrounding the specific error.

Example Snippets

U1: We need to consider factors A and B for making a decision in case X.

U2: Factor C would also be interesting and important to consider along with A and B.

Forgot Speaker	S: U2 mentions factor C will be important to take into consideration for case X.
Forgot Addressee	S: Thanks for bringing factors A, B and C up for case X, U1.
Wrong speaker	S: U1 mentions factor C will be taken into consideration for case X.
Wrong addressee	S: Interesting insight on factor C U1.

Example Snippets (Contd)

Wrong thread response

U1: This football season has been going great!
U2: I agree, for most teams anyway. Which one is your favorite?
U3: I prefer soccer instead. Anyone here a soccer fan?
U4: I don't really pay much attention to sports. My main hobby is movies!
U5: Yeah, and Knives Out was a great one!

S: I agree U5! The Rams are doing so well this year!

Inappropriately timed initiative

U1: I love documentaries and it has been great seeing so many come out in recent years.
U2: They do seem informative. I'm particularly interested in performative documentaries, they seem more personal.
U3: I also enjoy performative documentaries, like Supersize Me. Have you watched it U2?

S: Does anyone here like fiction?

Survey of existing literature

- Surveyed evaluation metrics utilized in past MPC modeling research tackling the tasks:
  - Speaker Identification
  - Response Selection or Generation
  - Addressee Recognition
  - ~15 papers total

Survey findings

- Most common metrics reported
  - BLEU
  - ROUGE
  - Classification reports
  - Yet most are reported on different properties (ex *n* is different for *n*-gram comparisons)
- We report inconsistencies across all literature

Need for better error reporting

- Most metrics reported are not consistent across the main task they focus
  - Even when reporting on shared task (DSTC-8 Track 2 NOESIS challenge)
  - Cannot compare across SOTA claims
- Not all models are publicly released
  - Difficult to re-evaluate even with possible new benchmarks

Next steps

- Formalize errors towards MPC modeling benchmark
  - Introduce automatic evaluation metrics
    - Classification reports for Speaker Identification and Addressee Recognition
    - Track interactions between the group
    - Graph similarity for conversation structure and thread management

Next steps (Contd)

- Formalize errors towards MPC modeling benchmark
  - Introduce human evaluation metrics
    - Naturalness
    - Belonging
    - Engagement
    - Initiative
    - + Towards all participants

Thank you!

References

[1] Sai, A. B., Mohankumar, A. K., & Khapra, M. M. (2022). A survey of evaluation metrics used for NLG systems. *ACM Computing Surveys (CSUR)*, *55*(2), 1-39.

[2] Celikyilmaz, A., Clark, E., & Gao, J. (2020). Evaluation of text generation: A survey. *arXiv preprint arXiv:2006.14799*.

[3] Deng, M., Tan, B., Liu, Z., Xing, E., & Hu, Z. (2021, November). Compression, Transduction, and Creation: A Unified Framework for Evaluating Natural Language Generation. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing* (pp. 7580-7605).

[4] Higashinaka, R., Araki, M., Tsukahara, H., & Mizukami, M. (2021, July). Integrated taxonomy of errors in chat-oriented dialogue systems. In *Proceedings of the 22nd Annual Meeting of the Special Interest Group on Discourse and Dialogue* (pp. 89-98).

[5] Novikova, J., Dušek, O., Curry, A. C., & Rieser, V. (2017, September). Why We Need New Evaluation Metrics for NLG. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing* (pp. 2241-2252).

[6] Van Der Lee, C., Gatt, A., Van Miltenburg, E., Wubben, S., & Krahmer, E. (2019). Best practices for the human evaluation of automatically generated text. In *Proceedings of the 12th International Conference on Natural Language Generation*(pp. 355-368).

[7] Howcroft, D. M., Belz, A., Clinciu, M. A., Gkatzia, D., Hasan, S. A., Mahamood, S., ... & Rieser, V. (2020, December). Twenty years of confusion in human evaluation: NLG needs evaluation sheets and standardised definitions. In *Proceedings of the 13th International Conference on Natural Language Generation* (pp. 169-182).

[8] Belz, A., Mille, S., & Howcroft, D. M. (2020, December). Disentangling the properties of human evaluation methods: A classification system to support comparability, meta-evaluation and reproducibility testing. In *Proceedings of the 13th International Conference on Natural Language Generation*(pp. 183-194).

[9] van Miltenburg, E., Clinciu, M. A., Dusek, O., Gkatzia, D., Inglis, S., Leppänen, L., ... & Wen, L. (2021, January). Underreporting of errors in NLG output, and what to do about it. In *INLG*.